

K-Bitcoin Script: An Executable Semantics for Modern Bitcoin Script

Author Name(s)

Department or Course Name

Institution Name

email@example.com

Abstract—Bitcoin Script is the consensus-critical stack language used to authorize transaction inputs. Although intentionally small, its operational behavior is distributed across byte-level script decoding, stack evaluation, soft-fork flags, transaction-dependent signature hashing, witness programs, and Taproot/Tapscript resource rules. This report presents K-Bitcoin Script, an executable semantics of Bitcoin Script in the K Framework. Following the style of prior large K semantics such as KEVM, KJS, and K-Java, we emphasize the engineering choices needed to make a language definition simultaneously executable, testable, and useful as a reference artifact. The semantics covers legacy, SegWit v0, and Taproot/Tapscript validation paths, including BIP-143, BIP-341, and BIP-342 signature hashing. We validate it against Bitcoin Core’s script and transaction test vectors and against a mainnet-derived benchmark dataset. On a 288,190-input dataset, the latest benchmark run produces 288,045 paired K/Core measurements with zero mismatches; K verifies the paired inputs in 175.4 seconds, compared with 84.2 seconds for libbitcoinconsensus.

Index Terms—Bitcoin Script, formal semantics, K Framework, Taproot, executable semantics, benchmarking

I. INTRODUCTION

Bitcoin consensus depends on faithful validation of each non-coinbase transaction input. For common transactions this validation is fast enough to look routine, but semantically it is dense: the interpreter must decode minimally encoded byte strings, maintain multiple stacks, enforce soft-fork flags, compute signature hashes over transaction context, switch execution phases for P2SH and witness programs, and implement distinct rules for legacy, SegWit, and Taproot-era spends. Production implementations such as Bitcoin Core are authoritative engineering artifacts, but they are not designed to be mathematical specifications.

The executable-semantics approach taken by the K Framework has been applied to large real-world languages and machines. KEVM formalizes the Ethereum Virtual Machine and evaluates both correctness and performance against the official Ethereum test suite [1]. KJS gives an executable semantics for JavaScript and uses the official ECMAScript conformance suite to test coverage [2]. K-Java formalizes Java 1.4 and reports a test-driven development process with a large conformance suite [3]. These papers share a useful pattern: make the semantics executable, define the configuration and transition rules in K, test the result against external conformance artifacts, and then demonstrate at least one concrete use of the semantics.

This project follows that pattern for Bitcoin Script. The goal is not to replace Bitcoin Core, but to build an independent artifact that is precise enough to run real inputs, readable enough to audit, and practical enough to use in regression testing and offline validation.

A. Motivations

Bitcoin Script is a good target for executable semantics for three reasons. First, it is consensus-critical: an implementation disagreement can split the network. Second, the language is bounded and deterministic for a fixed transaction context, making it more tractable than general-purpose languages. Third, Script’s historical evolution has produced many small semantic regimes: pre-P2SH, P2SH, strict DER signatures, locktime and sequence checks, SegWit v0, and Taproot/Tapscript. A formal model can expose the phase boundaries and flag-dependent behavior that are otherwise interleaved in implementation code.

B. Challenges

The central challenge is that script verification is not just stack-machine execution. Signature opcodes require transaction-dependent preimages; those preimages vary across legacy, BIP-143, BIP-341, and BIP-342 paths. Witness programs alter the script being executed and the data committed to by signatures. Tapscript changes opcode validity, introduces `OP_CHECKSIGADD`, and accounts for a signature validation budget.

The second challenge is performance. A semantics that can execute one test case is not automatically useful on mainnet-sized datasets. In early honest benchmarks, K spent most of its time moving large transaction byte strings through textual KORE and reconstructing large result patterns. The final benchmark in this report relies on binary KORE construction, subprocess chunking, native cryptographic hooks, and a done-phase cleanup rule that removes large no-longer-observed cells from final configurations.

C. Contributions

This report makes the following contributions.

- It describes an executable K semantics for modern Bitcoin Script, including legacy, SegWit v0, and Taproot/Tapscript validation phases.

- It explains how transaction context, previous outputs, witness data, signature hash modes, and consensus flags are represented in the K configuration.
- It reports validation against Bitcoin Core test vectors and a fresh mainnet benchmark run on the merged main branch.
- It documents the optimization path that reduced the honest K/Core performance ratio from 149× in the first complete run to 2.1× on the latest run.

II. BACKGROUND

A. Bitcoin Script

Bitcoin Script is a stack-based language used to specify spending conditions. A transaction input provides unlocking data, traditionally a `scriptSig`; the previous output provides the locking `scriptPubKey`. Validation executes these programs under consensus flags determined by block height and transaction context. P2SH adds a redeem script phase. SegWit v0 moves witness data outside the legacy transaction ID and changes the signature hash algorithm via BIP-143. Taproot introduces SegWit version 1 outputs, key-path spends, script-path spends, Schnorr signatures, and Tapscript validation rules via BIP-341 and BIP-342.

B. The K Framework

K definitions describe syntax, configurations, and rewrite rules. From a definition, K can generate parsers, interpreters, symbolic execution tools, and proof infrastructure. The prior K papers use this generated-tool approach as a methodological constraint: a semantics should be executable, tested against external programs, and reusable for more than prose documentation. K-Bitcoin Script adopts the same constraint. The repository contains about 5.2 KLOC of K definition files across modules for decoding, flow control, stack behavior, arithmetic, cryptography, signature checks, transaction parsing, and sighash construction.

III. K SEMANTICS OF BITCOIN SCRIPT

This section gives an implementation-oriented overview of the K definition. The goal is to make the report readable even to a reader who has not opened the repository. The definition follows the same discipline used by KEVM, KJS, and K-Java: the prose explains the organization, but the executable artifact is the semantics itself. In this project, that artifact is a collection of K modules under `script-semantics/`, backed by native hooks for cryptographic primitives and by a Python runtime wrapper that constructs the initial K configuration.

A. Definition Organization

The executable entry point is `script.k`. It imports a core semantics module and the modules that define decoding, flow control, stack operations, arithmetic, and signature behavior. The core module, `script-semantics.k`, owns the top-level configuration, phase sort, failure mechanism, phase transition rules, witness dispatch, and Taproot verification

skeleton. Supporting modules then define the local rewrite rules for particular opcode families.

Figure 1 shows the dependency shape. The important point is that `script-decode.k` is not merely a parser. It is part of the operational semantics: decoding emits K computations such as `#guardExec(OP_CHECKSIG)` and then recursively schedules decoding of the remaining byte string. This makes byte-level consensus checks, such as push-size limits and Tapscript `OP_SUCCESSx`, visible in the same rewrite system as ordinary stack-machine rules.

Table I summarizes the role of each major module. This split is useful for audit. A reader investigating stack behavior can stay mostly in `script-stack.k`; a reader investigating BIP-341 serialization can go to `script-sighash.k`; and a reader investigating phase boundaries can start in `script-semantics.k`. The main exception is `OP_CODESEPARATOR`, whose behavior necessarily crosses modules: the decoder records the byte tail after the separator, the flow and signature rules respect whether the opcode was actually executed, and the sighash module consumes the resulting script-code state.

B. Configuration

The top-level configuration is a single `<T>` cell containing the active computation, stacks, validation phase, error state, input scripts, witness blob, consensus flags, transaction context, and phase-specific signature metadata. Figure 2 groups the most important cells by role. The actual K configuration is slightly more detailed, but the grouping captures the design: ordinary opcodes see a small local state, while signature and phase rules can mention the extra cells needed for consensus validation.

The `<k>` cell starts at `#init`. Initialization checks `scriptSig` size and push-only requirements, then schedules decoding of either the unlocking script or the locking script. The data stack is `<stack>`; `<altstack>` models the Script altstack; and `<exec>` stores a list of booleans representing nested conditional-execution state. Most stack, arithmetic, and hash opcodes rewrite only `<k>` and `<stack>`. For example, a simple stack operation does not need to know the spending transaction, the previous output set, or the witness.

The phase and consensus cells make the same opcode name meaningful in the right validation regime. `<phase>` ranges over `scriptSig`, `scriptPubKey`, `p2sh-redeem`, `witness-v0`, `witness-v1`, and `done`. The `<flags>` bitmask models Bitcoin Core verification flags such as `FLAG_P2SH`, `FLAG_WITNESS`, `FLAG_MINIMALDATA`, and `FLAG_TAPROOT`. Rules can therefore express statements such as: “in witness phases, if `MINIMALIF` is active, `OP_IF` may only consume an empty vector or `0x01`.” This avoids baking one global interpretation of an opcode into the syntax.

The transaction context is explicit rather than hidden in a host callback. `<tx>` stores the serialized spending transaction, `<prevouts>` stores the previous outputs needed by SegWit and Taproot hashing, `<inputIndex>` selects the




Fig. 1. High-level organization of the executable K definition. The core configuration and phase machinery live in `script- semantics.k`; opcode families are separated into smaller modules, while cryptographic primitives are delegated to native hooks.

TABLE I
MAJOR K MODULES AND THEIR SEMANTIC RESPONSIBILITIES

Module	Primary content	Why it is separated
<code>script- semantics.k</code>	Configuration, phase sort, failure mechanism, witness dispatch, Taproot skeleton, cleanup rules	Owens global state and validation sequencing; most other modules import the declarations introduced here.
<code>script-decode.k</code>	Byte-to-opcode dispatch, push-data decoding, opcode-count accounting, Tapscript <code>OP_SUCCESSx</code> cases	Bitcoin Script is serialized bytecode; consensus behavior depends on how bytes are decoded, not just on the abstract opcode names.
<code>script-flow.k</code>	<code>OP_IF</code> , <code>OP_NOTIF</code> , <code>OP_ELSE</code> , <code>OP_ENDIF</code> , execution guarding, CLTV/CSV NOP behavior	Conditional execution determines whether most opcode rules fire or are skipped; centralizing the guard avoids duplicating branch logic.
<code>script-stack.k</code> , <code>script-arith.k</code> <code>script-sig.k</code>	Local stack and numeric opcode rules	These rules are mostly phase-independent and mention only the stack, altstack, or script-number helpers.
<code>script-sighash.k</code> , <code>script-tx.k</code> <code>script-crypto.k</code>	ECDSA and Schnorr signature opcode rules, multisig collection, <code>OP_CHECKSIGADD</code> Transaction field extraction, legacy/BIP-143/BIP-341/BIP-342 preimage construction Hash and signature hook boundaries	Signature opcodes are where stack state, consensus flags, transaction context, and phase-specific sighash construction meet. Sighash construction is byte-precise and large enough to deserve its own testable layer. Primitive cryptography is delegated to native code, while selection and serialization remain in K.

input under validation, and `<amount>` carries the spent-output amount. Older fixtures can still supply a precomputed `<sighash>` blob, but the main benchmark path uses K-side sighash construction from `<tx>` and `<prevouts>`. This is the difference between merely replaying Bitcoin Core’s answer and specifying how the answer is obtained.

Several cells exist specifically for script-code tracking. Legacy, BIP-143, and BIP-342 all depend on the portion of the script committed to by a signature, but they do so differently. The `<scriptCode>` cell stores the current BIP-143 or legacy script code. The `<pendingCodeSepTail>` cell is written by the decoder just before a decoded `OP_CODESEPARATOR`; if the separator is actually executed, the opcode rule can promote that tail into `<scriptCode>`.

TABLE II
SELECTED CONFIGURATION CELLS

Cell	Kind of state	Read by
<code><k></code>	Pending computation	All rules
<code><stack></code> , <code><altstack></code>	Script data stacks	Opcode rules
<code><exec></code>	Conditional branch stack	Flow guard, IF/ELSE/ENDIF
<code><phase></code>	Current validation phase	Phase, signature, decoder rules
<code><flags></code>	Consensus verification flags	Phase and opcode guards
<code><tx></code> , <code><prevouts></code>	Serialized tx context	Sighash rules
<code><witness></code>	Witness stack blob	Witness and Taproot dispatch
<code><scriptCode></code>	Active signature script code	Signature rules
<code><tapleafHash></code>	Tapscript leaf commitment	BIP-342 sighash
<code><error></code>	Bitcoin Core-style error code	Harness and tests

Tapscript uses `<tapleafHash>`, `<tapscriptBytes>`, and `<codeSepPos>` for the BIP-342 extension fields.

The configuration also includes a concrete-runtime cleanup




Fig. 2. Simplified view of the K configuration. Rewrite rules mention only the cells they need, which keeps ordinary opcode rules local while still allowing signature and phase rules to inspect transaction context.

rule. Once a successful run reaches `done` and the `<k>` cell drains, the LLVM backend rewrites large byte and list cells such as `<tx>`, `<prevouts>`, `<witness>`, and `<sighash>` to empty values. This rule is tagged concrete so proof-oriented backends do not see it. The rule does not change the result observed by the verifier, which reads only success/failure and the final error state, but it dramatically reduces the final KORE pattern that the runtime must serialize after each benchmark input.

C. Execution Phases

Script validation is modeled as a phase machine rather than as one monolithic interpreter loop. Figure 3 gives the high-level flow. Every ordinary script phase has the same inner shape: decode a byte string, execute the scheduled opcode computations, and then run `#phaseComplete` to decide the next phase. That decision depends on the stack, flags, script templates, witness data, and phase-specific validity checks.

The initial phase is `scriptSig`. If the unlocking script is empty, the semantics skips directly to `scriptPubKey`; otherwise it checks the 10,000-byte script-size limit and, when the flag is active, verifies that the unlocking script is push-only. After `scriptSig` completes, the rules clear the `altstack`, reset the opcode count, reset code-separator state, and schedule the locking script.

The `scriptPubKey` completion rules are the first major branch point. If the locking script is a P2SH template under `FLAG_P2SH`, the semantics restores the saved stack from `scriptSig`, extracts the redeem script, and enters `p2sh-redeem`. If it is a native witness program under `FLAG_WITNESS`, it dispatches to `#witnessVerify`. Otherwise the phase can finish directly, subject to truthiness, clean-stack, and unexpected-witness checks.

TABLE III
VALIDATION PHASES

Phase	Input executed	Main exit decision
<code>scriptSig</code>	Unlocking script	Schedule <code>scriptPubKey</code>
<code>scriptPubKey</code>	Locking script	Finish, P2SH, or witness dispatch
<code>p2sh-redeem</code>	Redeem script	Finish or witness dispatch
<code>witness-v0</code>	P2WPKH/P2WSH script	BIP-143 clean-stack completion
<code>witness-v1</code>	Tapscript	BIP-342 completion and budget rules
<code>done</code>	None	Runtime cleanup / observable success

The `p2sh-redeem` phase mirrors the post-`scriptPubKey` decision, but with the redeem script as the candidate witness program. This is how P2SH-wrapped SegWit spends are modeled. A redeem script that is not a witness program finishes through the ordinary clean-stack path; a redeem script that is a witness program enters the witness verifier. This split is one reason the phase cell is useful: P2SH redeem execution, native witness execution, and Tapscript execution all run decoded opcodes, but their boundary conditions differ.

Witness version 0 is split by program length. A 20-byte program is P2WPKH: the witness must contain exactly a signature and public key, and the semantics synthesizes the standard P2PKH script template before execution. A 32-byte program is P2WSH: the last witness item is the witness script, its SHA-256 hash must equal the program, and all preceding witness items become the initial stack. In both cases the phase becomes `witness-v0`, the witness is marked consumed, and `#witnessComplete` enforces exactly one truthy final stack item.

Witness version 1 dispatches to Taproot when `FLAG_TAPROOT` is set and the witness program is 32 bytes. A single effective witness item is a key-path spend: the semantics checks Schnorr signature shape, computes the BIP-341 sighash, and verifies the signature against the output key. More than one effective witness item is a script-path spend: the semantics validates the control block, computes the TapLeaf commitment, pushes the script-path witness items, decodes the tapscript, and enters `witness-v1`. Tapscript then finishes through `#tapscriptComplete`, again requiring one truthy item on the final stack.

Failures are explicit. The K item `#fail(ERR)` writes a Bitcoin Core-style error string into `<error>` and clears `<k>`. Negative test vectors are therefore not merely stuck terms: they terminate with structured failure labels such as `SCRIPT_SIZE`, `SIG_PUSHONLY`, `WITNESS_PROGRAM_MISMATCH`, `CLEANSTACK`, or `SCHNORR_SIG`. This is valuable for debugging because a mismatch can be classified as the wrong success value, the wrong error class, or a harness problem before looking at the full configuration trace.

D. Decoder and Guarded Execution

The decoder is the part of the definition that most directly reflects Bitcoin Script’s bytecode nature. It consumes a `Bytes` value, inspects the first byte, schedules the corresponding opcode computation, and recursively continues on the tail.




Fig. 3. Simplified phase machine for script verification. The curved #init-to-scriptPubKey edge is the empty-scriptSig path; the witness branch targets either v0 or v1 execution. The actual K rules include additional side conditions such as push-only checks, WITNESS_MALLEATED, WITNESS_UNEXPECTED, clean-stack checks, witness-program length checks, and Taproot control-block validation.

Figure 4 shows the pipeline. The decoder has to do more than map bytes to names: it enforces script-element size, counts opcodes above OP_16, identifies disabled and reserved opcodes, handles OP_SUCCESSx in Tapscript, and records byte tails for OP_CODESEPARATOR.

The central scheduling idiom is:

```
#guardExec(OP) ~> #decodeScript(rest).
```

The first component executes the opcode if the current conditional-execution stack is active; the second component continues decoding the remaining byte string. Flow-control opcodes such as OP_IF, OP_ELSE, and OP_ENDIF are emitted directly because they must update <exec> even in contexts where ordinary opcodes are skipped. This matches Bitcoin Script’s rule that disabled branches still have to be structurally balanced.

#guardExec is defined in script-flow.k. If every entry in <exec> is true, the guard rewrites to the opcode and the opcode rule fires. If any entry is false, the guard rewrites to .K, dropping the opcode without touching the stack. This small abstraction is one of the cleaner parts of the definition: it keeps skipped-branch behavior out of the hundreds of ordinary opcode rules. A stack opcode such as OP_DUP can be written as a local stack rewrite; the guard is responsible for deciding whether that local rewrite is even reached.

Push-data decoding is also semantic. OP_PUSHBYTES_1 through OP_PUSHBYTES_75, OP_PUSHDATA1, OP_PUSHDATA2, and OP_PUSHDATA4 all compute a length, extract that many bytes, enforce the 520-byte push-size limit, and emit OP_PUSHDATA(opcode, data). Minimal-push checks are intentionally performed at execution time instead of pure decode time, because MINIMALDATA should not reject a non-minimal push inside a branch that is not executed. This is a good example of where an executable semantics has to follow the consensus interpreter’s control-flow sensitivity rather than a cleaner context-free parser design.

The decoder is also phase-sensitive in Tapscript. Several byte values that are invalid or reserved in earlier phases become OP_SUCCESSx in witness-v1. The K rule for such a byte checks <phase>; in Tapscript it schedules #opSuccess, which clears the remaining computation and

TABLE IV
DECODER RESPONSIBILITIES

Responsibility	K mechanism
Script byte iteration	#decodeScript and #decode1
Opcode resource limit	#countOp, #countOps
Push extraction	#decodePush(opcode, length, rest)
Conditional skipping	#guardExec(OP) over <exec>
Minimal-push timing	Execution-time #isMinimalPush checks
Code separator tracking	#setCodesepTail before opcode execution
Tapscript success opcodes	Phase-sensitive #opSuccess rules

leaves a truthy value on the stack. Outside Tapscript, the same byte may become a reserved-opcode failure or an invalid-opcode failure. This is another reason phase is explicit state rather than a parameter hidden in the harness.

E. Opcode Families

After decoding, most opcodes fall into small rule families. Stack opcodes rewrite <stack> and occasionally <altstack>. Arithmetic opcodes convert byte vectors into Bitcoin Script numbers, enforce numeric-size limits, perform an integer operation, and push the minimally encoded result. Hash opcodes pop one byte vector, call a hash hook, and push the resulting digest. Locktime and sequence opcodes consult <nLockTime>, <nSequence>, and verification flags.

Signature opcodes are the exception to the local-rule pattern. Legacy OP_CHECKSIG must know the public key, signature, hash type, active script code, and transaction serialization. OP_CHECKMULTISIG has to collect variable-length public-key and signature lists, handle the historical dummy element, count signature operations, and apply NULLDUMMY/NULLFAIL. Tapscript adds OP_CHECKSIGADD, Schnorr signature-shape checks, public-key-type rules, and a signature validation budget. For this reason the signature module is where many cells meet: <stack>, <phase>, <flags>, <scriptCode>, <tx>, <prevouts>, <inputIndex>, <amount>, <tapleafHash>, and <sigopsWeight>.

This organization keeps the readable part of the semantics close to the language structure. A reader can treat ordinary opcodes as small rewrite rules and signature opcodes as phase-dispatching rules. The latter select between legacy, BIP-143, BIP-341, and BIP-342 sighash helpers, then call the




Fig. 4. Decoder pipeline. The decoder schedules executable K items while preserving byte-level resource checks and phase-sensitive special cases.

appropriate ECDSA or Schnorr hook. The hook boundary is deliberately narrow: the hook verifies a cryptographic predicate over bytes, while K constructs the bytes and chooses the predicate.

F. Signature Hashing

Signature checking is the largest semantic surface in the project. Legacy ECDSA checks use the legacy preimage rules, with support for historical hash-type behavior. SegWit v0 checks route through BIP-143. Taproot key-path and Tapscript checks route through BIP-341 and BIP-342. Native hooks implement primitive cryptographic operations and bulk SHA-256 walkers, while the K rules own the selection logic, serialization layout, and phase-dependent dispatch.

This boundary follows the precedent of prior K definitions: primitive operations may be hooked for execution, but semantic structure stays in K. The result is a definition that can run mainnet-sized transactions without asking K to perform every byte-level hashing loop one byte at a time. It also makes the validation story more meaningful. When K and Core agree on a Taproot input, the agreement covers the witness layout, annex handling, control-block validation,

TapLeaf hash, sighash serialization, and Schnorr predicate, not just an externally supplied digest.

Figure 5 shows the dispatch at a high level. Legacy and P2SH execution use ECDSA over the historical transaction digest. SegWit v0 uses ECDSA over BIP-143, with `<scriptCode>` set either to the canonical P2WPKH template or to the current P2WSH witness script. Taproot key-path spends use BIP-341 directly. Tapscript spends use BIP-342’s extension fields: the tapleaf hash, key version, and code-separator position are part of the signed message. The K definition keeps these choices in the rewrite rules rather than in the wrapper, so a phase bug appears as a semantic bug, not as an invisible host-language dispatch.

G. Spend-Shape Walkthroughs

The phase machine is easiest to read through common spend shapes. Table V lists paths for representative cases. These are not additional algorithms beyond the K rules; they are useful mental traces through the executable definition.

Consider native P2WPKH as a compact example. The initial `scriptSig` is empty, so `#init` schedules the locking script directly and sets the phase to `scriptPubKey`. The locking script decodes to a witness program. On phase




Fig. 5. Signature-hash dispatch. The signature opcode rules select a preimage algorithm from the current phase and then call a narrow cryptographic hook over the bytes constructed by K.

TABLE V
COMMON SPEND SHAPES AS PATHS THROUGH THE K SEMANTICS

Spend shape	Phase path	Key semantic events
Legacy P2PKH or bare script P2SH	scriptSig → scriptPubKey → done scriptSig → scriptPubKey → p2sh-redeem → done	scriptSig pushes unlocking data; scriptPubKey consumes it; signature opcodes use the legacy preimage and ECDSA hook. The stack after scriptSig is saved; the redeem script is popped from the saved stack and decoded only after the P2SH template succeeds.
Native P2WPKH	scriptPubKey → witness-v0 → done	Empty scriptSig; witness must contain signature and pubkey; K synthesizes the standard P2PKH template and uses it as BIP-143 <scriptCode>.
Native P2WSH	scriptPubKey → witness-v0 → done	Last witness item is hashed and compared with the witness program; earlier witness items become the initial stack; the witness script becomes <scriptCode>.
Taproot key path	scriptPubKey → #taprootVerify → done	A single effective witness item is checked as a Schnorr signature over the BIP-341 message; no tapscript is decoded.
Taproot script path	scriptPubKey → #taprootVerify → witness-v1 → done	Control block validates the output key; K computes the TapLeaf hash, pushes script-path witness items, decodes the tapscript, and enforces BIP-342 completion rules.

completion, the rules check that the witness flag is active and that the unlocking script was empty, then call `#witnessVerify(0, PROG)`. Because `PROG` is 20 bytes, the rules require exactly two witness items, push them onto `<stack>`, build the canonical P2PKH byte string, store that byte string in `<scriptCode>`, and decode it under `witness-v0`. The ordinary P2PKH opcodes then run through the same stack and signature rules used elsewhere, but `OP_CHECKSIG` dispatches to BIP-143 because the phase is `witness-v0`. Finally, `#witnessComplete` requires exactly one truthy stack item.

P2WSH follows the same `witness-v0` completion discipline but differs in how the executing script is chosen. The last witness item is a byte string, not a template synthesized by the semantics. K computes its SHA-256 hash and requires equality with the 32-byte witness program. The earlier witness items become the stack, and the witness script becomes both the decoded program and the initial BIP-143 `<scriptCode>`. `OP_CODESEPARATOR` can then update `<scriptCode>` during execution, but only when the separator is in an executed branch.

Taproot script path is the richest trace. `#taprootVerify` first removes a possible annex from the effective witness count, validates the control-block size, computes the merkle path, and checks that the tweaked key matches the output key. Only after those structural checks succeed does the semantics enter

`witness-v1`. At that point the tapscript bytes, tapleaf hash, and code-separator position cells have been initialized for BIP-342. Tapscript opcodes then execute under the decoder's phase-sensitive rules, which is what makes `OP_SUCCESSx` successful in this phase but not in earlier phases.

H. Runtime Interface

The Python wrapper constructs initial configurations and invokes the compiled LLVM backend. The latest runtime path avoids serializing the whole initial configuration through textual KORE. Instead, the wrapper uses the K LLVM C API to build tokens and injections directly and calls the binary runtime path from `verify_script`. The benchmark runner also goes through `verify_script`, ensuring the measured path is the same optimized path used by normal verification.

The runtime interface is intentionally thin. It translates host data into K cells, starts the semantics, and translates the final success or error state back to Python. Consensus behavior is not encoded in Python branches except for preparing the input material. This matters for the report's claim: the artifact being evaluated is the K semantics, not a Python reimplementaion that happens to call K for a subset of opcodes.

IV. VALIDATION

A. Reference Tests

Table VI summarizes the current validation story. The semantics passes all 1,222 Bitcoin Core

script_tests.json vectors. It also passes 201 of 214 transaction vectors; the remaining 13 are documented expected failures, consisting of malformed BADTX cases and legacy CONST_SCRIPTCODE sighash edge cases.

TABLE VI
REFERENCE TEST COVERAGE

Suite	Cases	Passing	Notes
script_tests.json	1,222	1,222	Standard, SegWit, Taproot
tx_valid.json	121	121	Transaction validation
tx_invalid.json	93	80	13 expected xfails
Taproot vectors	26	26	Key path, control blocks, annex

B. Sighash Tests

The project includes direct tests for legacy, BIP-143, and BIP-341 sighash entry points, plus C-hook isolation tests that compare hook output with Python `hashlib` implementations. BIP-342 sighash is currently covered indirectly through full script verification on mainnet Tapscript inputs rather than by a standalone published vector file. This is a useful distinction: mainnet verification is strong end-to-end evidence, but a direct BIP-342 vector suite would make the coverage argument cleaner.

C. Mainnet Dataset

The benchmark dataset contains 288,190 real mainnet inputs from 72 blocks, spanning pre-P2SH through Taproot-era transactions. K is run in 29 subprocess chunks to bound the LLVM runtime arena size; each chunk drops five warmup inputs, producing 288,045 K result rows. Core is run as a separate Core30 pass and drops five warmup inputs once, producing 288,185 Core rows. The combined report contains the 288,045 paired rows for which both K and Core results are available. All paired rows agree.

V. QUANTITATIVE EVALUATION

A. Methodology

The final benchmark was run on the remote M4 Pro machine previously used for the v8-v10 benchmark series. The repository was synced to `origin/main` after PR #11, commit `d11942b`. The K LLVM definition was rebuilt from that commit. K was measured with one iteration per input using the chunked subprocess driver. Core was measured with `libbitcoinconsensus` from Bitcoin Core 27.2, using 30 iterations per input to match the prior v10 artifact.

B. Overall Results

Table VII gives the main result. K verifies the paired dataset in 175.4 seconds, while Core takes 84.2 seconds. The wall-time ratio is $2.1\times$.

The median K noop baseline is $234.1\ \mu\text{s}$, so the median net verification cost is about $141.8\ \mu\text{s}$. Core’s median noop baseline is about $1.0\ \mu\text{s}$, with a median net verification cost of about $26.1\ \mu\text{s}$.

TABLE VII
FRESH MERGED-MAIN BENCHMARK, PAIRED ROWS

Engine	Inputs	Total	Throughput	Median
K Framework	288,045	175.4 s	1,642.3/s	375.8 μs
libbitcoinconsensus	288,045	84.2 s	3,419/s	27.1 μs

C. Results by Era

Table VIII shows median and 95th-percentile latency by consensus era. DERSIG has a high Core median because the dataset includes large historical stress transactions. The K median remains under $400\ \mu\text{s}$ across all eras, while its DERSIG p95 records the residual cost of the stress cases.

TABLE VIII
LATENCY BY CONSENSUS ERA

Era	Inputs	K med	K p95	Core med
pre-P2SH	286	350.4 μs	387.3 μs	15.5 μs
P2SH	8,840	353.8 μs	484.5 μs	17.3 μs
DERSIG	59,479	364.0 μs	4.8 ms	2.1 ms
CLTV	42,167	385.8 μs	880.0 μs	27.5 μs
CSV	39,361	379.6 μs	834.1 μs	23.4 μs
SegWit	70,097	379.7 μs	1.1 ms	21.6 μs
Taproot	67,815	372.6 μs	821.3 μs	22.9 μs

D. Representative and Stress Inputs

The benchmark separates representative inputs from stress inputs. On representative inputs, K takes 105.3 seconds and Core takes 13.3 seconds, for a $7.9\times$ total-time ratio. On stress inputs, K takes 70.1 seconds and Core takes 70.9 seconds, putting the two essentially at parity. This result is important because the first honest K-native benchmark was dominated by stress transactions.

TABLE IX
BENCHMARK TOTALS BY CATEGORY

Category	Inputs	K total	Core total
Representative	232,884	105.3 s	13.3 s
Stress	55,161	70.1 s	70.9 s

E. Optimization History

Table X summarizes the performance path. The v4 run was fast but semantically misleading because it used a precomputed sighash shortcut for most inputs. The v8 run was the first complete honest K-native run and showed a $149\times$ gap. Binary KORE construction reduced uniform boundary overhead in v9. The v10 cleanup removed large cells from final result patterns and fixed an $O(n^2)$ previous-output blob construction in the benchmark runner.

TABLE X
BENCHMARK HISTORY

Run	K total	Core total	Ratio	Note
v4	198 s	84 s	2.4×	Sighash shortcut
v8	12,652 s	85 s	149×	Honest K-native
v9	7,388 s	85 s	87×	Binary FFI
merged main	175.4 s	84.2 s	2.1×	Cleanup + FFI

VI. APPLICATIONS

The immediate application is an executable reference model for testing consensus-sensitive behavior. Because the same K rules execute script tests, transaction vectors, and mainnet-derived inputs, regressions in opcode semantics or sighash dispatch can be caught without creating a second model. The benchmark harness also functions as a performance regression test: it identified when the v10 binary path was present in the code but bypassed by a later benchmark runner change.

The next natural application is offline audit. A K-backed verifier can replay selected blocks or transactions independently of Bitcoin Core’s interpreter and flag disagreements for investigation. This is not a node replacement; it is a high-assurance cross-checking tool.

VII. RELATED WORK

KEVM is the closest methodological predecessor. It formalizes a consensus-critical bytecode machine, evaluates against official tests, and then demonstrates verification applications for smart contracts [1]. K-Bitcoin Script shares the executable-semantics goal but targets a deliberately non-Turing-complete authorization language rather than a gas-metered virtual machine.

KJS demonstrates that a large, messy production language can be specified in K and validated against an external conformance suite [2]. K-Java makes a similar argument for Java and emphasizes test-driven semantic development and a large suite of feature tests [3]. This project follows the same discipline: external tests and real programs are not afterthoughts, but part of the definition process.

Bitcoin Core remains the de facto executable reference for Bitcoin consensus. This project does not replace it. Instead, it provides a separately written, auditable semantics whose behavior can be compared with Core on identical inputs.

VIII. LIMITATIONS AND FUTURE WORK

The current semantics still has documented gaps. Thirteen Bitcoin Core transaction vectors are expected failures, and direct BIP-342 sighash vectors should be added to complement mainnet Tapscript coverage. Benchmark artifacts are large and currently live on the remote benchmark machine rather than in the repository; they should be published as release artifacts or archived with the final report. The performance numbers are single-run measurements. A final paper-quality evaluation should include repeated runs, variance, exact hardware details, and software versions.

Future performance work includes caching shared transaction configuration across inputs from the same transaction and parallelizing independent input verification. Future semantic work includes making the remaining transaction vector gaps explicit in the K rules or in the harness and expanding standalone vector coverage for Tapscript sighash behavior.

IX. CONCLUSION

K-Bitcoin Script shows that an executable K semantics can cover modern Bitcoin Script, including SegWit and Taproot-era behavior, while remaining practical enough to run hundreds of thousands of real mainnet inputs. The latest merged-main benchmark reports zero mismatches over 288,045 paired K/Core measurements and a 2.1× wall-time ratio against `libbitcoinconsensus`. The result supports the central thesis of the K formalization papers: a semantics can be both formal and operational, provided it is engineered, tested, and benchmarked as a real implementation.

REFERENCES

- [1] E. Hildenbrandt, M. Saxena, N. Rodrigues, X. Zhu, P. Daian, D. Guth, B. Moore, D. Park, Y. Zhang, A. Stefanescu, and G. Rosu, “KEVM: A Complete Formal Semantics of the Ethereum Virtual Machine,” in *IEEE Computer Security Foundations Symposium*, 2018, doi: 10.1109/CSF.2018.00022.
- [2] D. Park, A. Stefanescu, and G. Rosu, “KJS: A Complete Formal Semantics of JavaScript,” in *ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2015, doi: 10.1145/2737924.2737991.
- [3] D. Bogdanas and G. Rosu, “K-Java: A Complete Semantics of Java,” in *ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, 2015, doi: 10.1145/2676726.2676982.

- [4] K Framework contributors, "The K Framework." [Online]. Available: <https://kframework.org/>
- [5] Bitcoin Core contributors, "Bitcoin Core." [Online]. Available: <https://bitcoincore.org/>
- [6] P. Wuille and J. Lau, "BIP 143: Transaction Signature Verification for Version 0 Witness Program," 2016. [Online]. Available: <https://github.com/bitcoin/bips/blob/master/bip-0143.mediawiki>
- [7] A. Poelstra, P. Wuille, J. Nick, and T. Ruffing, "BIP 341: Taproot: SegWit version 1 spending rules," 2020. [Online]. Available: <https://github.com/bitcoin/bips/blob/master/bip-0341.mediawiki>
- [8] A. Poelstra, P. Wuille, J. Nick, and T. Ruffing, "BIP 342: Validation of Taproot Scripts," 2020. [Online]. Available: <https://github.com/bitcoin/bips/blob/master/bip-0342.mediawiki>